

CS 3113 Introduction to Operating Systems


Topic #3. Processes

Outline

- 3.5 IPC in Shared-Memory Systems
- 3.6 IPC in Message-Passing Systems
- 3.7 Examples pf IPC Systems
- 3.8 Communication in Client–Server Systems


Communication Models

Shared Memory



(a)

Message Passing



(b)

IPC – Shared Memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of these processes & not the operating system.
 - Good for efficiency (no system calls to update memory)
 - Lots of opportunities for bugs
- Execution of program started via GUI mouse clicks, command line entry of its name, etc.
 - Resides in the address space of the process creating the shared memory
 - Other processes must attach the shared memory to their address space

IPC – Shared Memory (Cont'd)

Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.

Producer-Consumer Problem

- Producer: process generates data through some mechanism
- Consumer: process uses data generated by another

Producer-Consumer Problem

Typical approach: implement a data buffer from the producer to the consumer

- **unbounded-buffer** places no practical limit on the size of the buffer
- **bounded-buffer** assumes that there is a fixed buffer size

Circular/Shared Buffer of Items

- Items are instances of type **item**

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- in = the next free location to place a new item
- out = the next full location to remove an item from
- Both the producer and consumer processes have access to this buffer

Circular/Shared Buffer of Items

- $in == out$: no items in the buffer
- $(in+1) \% BUFFER_SIZE == out$: buffer is full

Circular Buffer: Producer

```
item next_produced;
while (true) {
    // Generate new item
    next_produced = ...

    // Wait for there to be space in the buffer
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */

    // Place item in the buffer
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Circular Buffer: Consumer

```
item next_consumed;
while (true) {
    // Wait for item to be available
    while (in == out)
        ; /* do nothing */

    // Get the next item
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    // Do something with the item
}
```

Outline

- 3.5 IPC in Shared-Memory Systems
- 3.6 IPC in Message-Passing Systems
- 3.7 Examples of IPC Systems
- 3.8 Communication in Client–Server Systems

IPC in Message Passing Systems

Mechanism for processes to communicate and to synchronize their actions

- Message system: processes communicate with each other without resorting to shared variables
- IPC facility provides two generic operations:
 - send(message)
 - receive(message)
- The message size can be either fixed or variable

Message Passing

- If processes P and Q wish to communicate, they need to:
 - Establish a ***communication link*** between them
 - Exchange messages via send/receive
- Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link? (buffer)
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?

Message Passing

Implementation of a communication link

- Physical choices:
 - Shared memory
 - Hardware bus
 - Network
- Logical choices:
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering

Direct Communication Model

- Processes must name each other explicitly:
 - **send** (*P, message*): send a message to process P
 - **receive**(*Q, message*): receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication Model

- Messages are directed to and received from **mailboxes** (also referred to as **ports**)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication Model

- Operations
 - create a new mailbox (port)
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - send**(*A, message*) – send a message to mailbox A
 - receive**(*A, message*) – receive a message from mailbox A

Mailbox Sharing

- *Scenario*
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?
- *Solutions*
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Message passing may be either blocking or non-blocking
- Blocking is considered synchronous
 - Blocking send: the sender is blocked until the message is received
 - Blocking receive: the receiver is blocked until a message is available

Synchronization

- Non-blocking is considered asynchronous
 - Non-blocking send: the sender sends the message and continues
 - The message is placed into a temporary buffer
 - Non-blocking receive: the receiver receives:
 - A valid message, or
 - Null message (nothing to receive)
- Different combinations possible
 - If both send and receive are blocking, we have a **rendezvous**

Synchronization with Rendezvous

Producer-consumer becomes trivial

```
message next_produced;  
while (true) {  
    /* produce an item in next produced */  
    send(next_produced);  
}
```

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next consumed */  
}
```

Buffering

- Queue of messages attached to the link
- Implemented in one of three ways:
 - Zero capacity: no messages are queued on a link.
Sender must wait for receiver (rendezvous)
 - Bounded capacity: finite length of n messages
 - Sender must wait if link full
 - Receiver must wait if link has nothing
 - This is our circular buffer example!
 - Unbounded capacity: infinite length
Sender never waits

Buffering (Cont'd)

Advantages:

- Processes don't have to wait for each other
- -> Potentially more efficient use of CPU and other resources

Challenges:

- Capacity must be large enough
- Consumer can potentially fall far behind
 - If timing is important, this can be a big issue

Quiz

- What are advantages and disadvantages of an unbounded capacity buffer?

Quiz answer

- **Advantages:** the operating-system provides the buffers in the form of a queue with indefinite length. This means that any number of messages can wait in the queue, so the sender will never have to block
- **The disadvantage** of automatic buffering is that it involves a more complex system that may be difficult to implement so that it manages memory efficiently. A scheme may reserve a sufficiently large memory space for the messages; however the memory may never fully be used, thereby wasting memory space.

Outline

- 3.5 IPC in Shared-Memory Systems
- 3.6 IPC in Message-Passing Systems
- 3.7 Examples of IPC Systems
- 3.8 Communication in Client–Server Systems

Shared Memory in POSIX (includes Linux)

- Process first creates shared memory segment (or opens an existing one):

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Set the size of the object:

```
ftruncate(shm_fd, 4096);
```

- Create a pointer to the shared memory:

```
shared_memory = mmap(shm_fd, 0, 4096 PROT_WRITE,  
MAP_SHARED, shm_fd, 0);
```

- Now the process can write to the shared memory

```
sprintf(shared_memory, "Writing to shared memory");
```

Shared Memory Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```

Shared Memory Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm.open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm.unlink(name);

    return 0;
}
```

Pipes

- Act as a conduit allowing two processes on the same computer to communicate
- Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., parent-child) between the communicating processes?
 - Can the pipes be used over a network?

Types of Pipes

- Ordinary pipes: cannot be accessed from outside the process that created it.
 - Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- Named pipes: can be accessed without a parent-child relationship.

Ordinary Pipes Notes

- Parent creates the process and then forks()
- Both parent and child close one half of the pipe so there is only one reader and one writer for the pipe
 - Either the parent or the child can become the producer
- When the producer closes the fd, the bytes in the buffer are still available for reading
 - When there are no more bytes, the consumer will receive an EOF marker

Named Pipes Notes

- One process opens for reading, the other for writing
- `open()` system call blocks until the other process also calls `open()`
- When the producer closes, the bytes in the pipe's buffer are still available for reading
 - When there are no more bytes, the consumer receives an EOF marker
- If you wrap a `FILE*` around a pipe fd, then make sure that you flush the `FILE`'s buffer (different than the pipe buffer) to make sure that the bytes are actually sent

Ordinary Pipes


Ordinary Pipes allow communication in standard producer-consumer style

- Producer writes to one end (the write-end of the pipe)
- Consumer reads from the other end (the read-end of the pipe)
- Ordinary pipes are therefore unidirectional

Ordinary Pipe

```
int fds[2];  
int ret = pipe(fds);  
if(ret < 0) exit(-1);  
/* Now use the pipe */
```

- `fds[0]`: output from pipe
- `fds[1]`: input to pipe



Ordinary Pipes for Communication

- The pipe is implemented inside the kernel (so, it does not exist within the process)
- However, the process maintains this pair of file descriptors, which allow it to reference the pipe


Ordinary Pipes for Communication

- The file descriptors cannot be shared outside of the process
- But: if the process forks(), then both the parent and child will have copies of the file descriptors!
 - And these reference the same pipe


Ordinary Pipe with Fork()

```
int fds[2];
int ret = pipe(fds);
if(ret < 0) exit(-1);
/* Now use the pipe */
int pid = fork()
if(pid > 0) {          /* Note: leaving off error case*/
    /* parent code */
} else {
    /* child code */
}
```

Before Fork



After Fork




Pipes

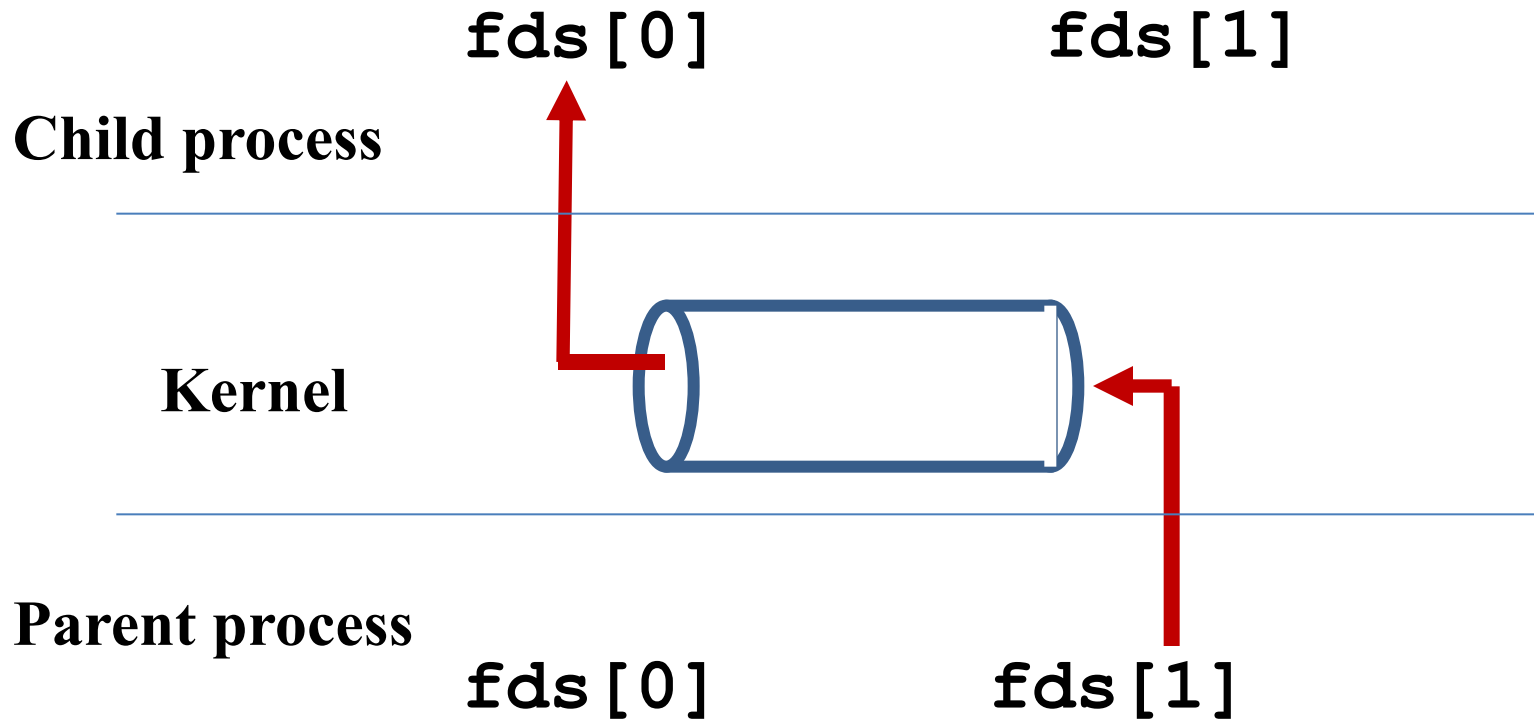
For our purposes, we will assume:

- Pipes will only work properly with one reader and one writer (this is the typical use case)
- This means, after the fork, each of the parent and child will close one of the two file descriptors

After Fork: Parent is the Producer



After Fork: Parent is the Producer



After Fork: Parent is the Producer

```
// Now use the pipe
int pid = fork()
if(pid > 0) {                                // Note: leaving off error case
    // parent code
    close(fds[0])
    :
}else {
    // child code
    close(fds[1])
    :
}
```

After Fork: Parent is the Producer

After fork and closing:

- Parent can write bytes to `fds[1]`
- Child can read these bytes from `fds[0]`

After Fork: Parent is the Producer

- The pipe has a buffer: it will hold written bytes until they are read
- If the writer closes the pipe, then
 - The reader gets to read the remaining bytes
 - But then will see an EOF
- Windows calls ordinary pipes ***anonymous pipes***

Named Pipes

Named Pipes are more powerful than ordinary pipes

- Communication *can* be bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems

Named Pipes in Unix

- Access points exist in the file system
 - Open them just as you would a file!
 - Use read()/write() to receive/send data
- Can have multiple readers/writers
 - A message is delivered to one randomly selected reader
 - So, effectively, they are bidirectional
 - **However, we will use them as unidirectional pipes**
- Create at the command line (or programmatically):
`mkfifo [NAME]`

Outline

- 3.5 IPC in Shared-Memory Systems
- 3.6 IPC in Message-Passing Systems
- 3.7 Examples of IPC Systems
- 3.8 Communication in Client-Server Systems

Client-Server Model of Communication


- Built on top of Producer/Consumer concept
- Server: provides some service (data, computation)
- Client: requests actions on the part of the server
- Implementation choices include:
 - Sockets
 - Pipes
 - Remote Procedure Calls
 - Remote Method Invocation (Java)

Sockets

A socket is defined as an endpoint for communication

- Identified by a concatenation of IP address and port: a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to **port 1625** on **host 161.25.19.8**
- All communication is done between a pair of sockets (one for client; the other for server)
- All ports below 1024 are well known and are used for standard services. Others are available for users
- Special IP address 127.0.0.1 (loopback) to refer to system on which process is running

Socket Communication



Remote Procedure Calls (RPCs)

- From the programmer's perspective, they appear as functions/methods that take arguments and return a value
- Under the hood, this function call:
 - Contacts a server
 - Sends the arguments to the server
 - Server does the work and sends the result back
 - Return the return value back to the client

Procedure for Using Shared Memory

- Find a *key*: Unix uses this key for identifying shared memory segments
- Use `shmget()` to allocate a shared memory
- Use `shmat()` to attach a shared memory to an address space
- Use `shmdt()` to detach a shared memory from an address space
- Use `shmctl()` to deallocate a shared memory

Project 1

➤ Header (incomplete)

```
#include <sys/ipc.h>
```

```
#include <sys/shm.h>
```

```
#include <unistd.h>
```

```
#include <errno.h>
```

```
/* key number */
```

```
#define SHMKEY ((key_t) 1497)
```


Key

- Unix uses this key for identifying shared memory segments.
 - A key is a value of type `key_t`
 - There are three ways to generate a key
 - Do it yourself
 - Use function `ftok()`
 - Ask the system to provide a private key

Key

- Do it yourself:
 - `#define SHMKEY ((key_t) 1497)`
- Use `ftok()` to generate one for you
 - `key_t = ftok(char *path, int ID);`
 - `path` is a path name (e.g., `“./”`)
 - `ID` is an integer (e.g., `‘a’`)
 - Function `ftok()` returns a key of type `key_t`: `Key = ftok(“./”, ‘x’)`
- Keys are global entities. If other processes know your key, they can access your shared memory.
- Ask the system to provide a private key using `IPC_PRIVATE`

Project 1

- shared memory

```
typedef struct
```

```
{
```

```
    int value;
```

```
} shared_mem;
```

```
shared_mem *total;
```

- process1 increases the value of shared variable "total" * by some number
- process2
- process3
- process4

Project 1

➤ Main function

```
int  shmid, pid1, pid2, pid3, pid4, ID, status;
```

```
char *shmadd;
```

```
shmadd = (char *) 0;
```

➤ `/* Create and connect to a shared memory segment*/`

```
if ((shmid = shmget (SHMKEY, sizeof(int), IPC_CREAT | 0666)) < 0){  
    perror ("shmget");  
    exit (1);};
```

```
if ((total = (shared_mem *) shmat (shmid, shmadd, 0)) ==  
(shared_mem *) -1) {  
    perror ("shmat");  
    exit (0);};
```

Project 1

- Initialize shared memory to 0

```
total->value = 0;
```

- Create processes

```
if ((pid1 = fork()) == 0)
```

```
    process1();
```

- Create additional processes

Project 1

- Parent wait for child processes to finish and print ID of each child. Three ways:
 - `wait(&status)`
 - `wait(NULL)`
 - `waitpid(pid1, NULL, 0);`
 - `printf("Child with pid %d has just exited.\n", pid1);`

Project 1

- To detach a shared memory, use `shmdt(total);`
 - `total` is the pointer returned by `shmat()`

```
if (shmdt(total) == -1) {  
    perror ("shmdt");  
    exit (-1);  
}
```
- To remove a shared memory, use `shmctl(shmid, IPC_RMID, NULL);`
 - `shmid` is the shared memory ID returned by `shmget()`.
- After a shared memory is removed, it no longer exists.

Notes

- If you did not remove your shared memory segments (e.g., program crashes before the execution of `shmctl()`), they will be in the system forever. This will degrade the system performance.
- Use the `ipcs` command to check if you have shared memory segments left in the system.
- Use the `ipcrm` command to remove your shared memory segments.